

# Cloud Computing Final Report

## 1 Introduction

Large-scale deep neural networks (DNNs) have achieved great success in various fields, such as computer vision (CV), natural language processing (NLP), speech recognition and so on. With the yearning for deep learning democratization (Garvey, 2018), there are increasing demands to deploy DNNs on resource-constrained devices (i.e., FPGA, GPU, etc.). However, there are many challenges to accommodate DNN models on hardware, such as limited storage, limited computational speed and high demand of real-time inference. Based on these challenges, model compression has been investigated in recent years. The core idea of model compression is to generate a much sparse model thus we could get acceleration in computation and reduction on space. Several model compression techniques were proposed to find a sparse network on DNNs to meet the demands, such as weight pruning (Han et al., 2015a), knowledge distillation (Hinton et al., 2015), quantization (Zhang et al., 2020), low rank approximation (Yu et al., 2017), sparsity regularization (Louizos et al., 2017), etc.

Among all the compression techniques referred above, weight pruning (Zhang et al., 2018a) is one of the most simple but popular one and is widely explored in recent years. The technique is implemented by zero out certain weights and then optimize the rest, which is a destruction plus learning process. From the aspect of hardware friendly or not, it can be divided into two categories which are structured pruning (hardware friendly) (Anwar et al., 2017; Wang et al., 2019) and unstructured pruning (hardware unfriendly). For structured pruning, we will explore pruning methods such as row pruning, column pruning and whole block pruning while for the unstructured pruning, we will conduct experiment on irregular pruning. In this work, we will demonstrate the effectiveness of weight pruning in the fields of both CV and NLP.

## 2 Background

Advances in Machine Learning have enabled high accuracy in classifications, recommendations, natural language processing, etc. The success of modern deep neural networks is mainly dependent on the availability of advanced computing power and a large number of data. (Wang et al., 2021) Technically, a large enough neural network and dataset could result in a machine that can do anything a human brain is capable of. However, a high-performance neural network usually means complicated and has a huge amount of parameters to update. The enormous size of a network not only causes great challenges on hardware devices but also wastes unnecessary computation and time consumption. A simple way to avoid such wastes is to prune the parameters in the network that have slight impacts.

### 2.1 Deep Neural Network

A deep neural network (DNN) is a framework for deep learning, which is a neural network with at least one hidden layer. Similar to shallow neural networks, deep neural networks can also provide modeling for complex nonlinear systems, but the extra levels provide a higher level of abstraction for the model, thus improving the capabilities of the model (Chen et al., 2019).

The deep neural network is a discriminative model that can be trained using a backpropagation algorithm. The weight update can be solved by the stochastic gradient descent method using the following formula (Chen et al., 2019):

$$\Delta\omega_{ij}(t+1) = \Delta\omega_{ij}(t) + \eta \frac{\partial C}{\partial \omega_{ij}} \quad (1)$$

Among them,  $\eta$  is the learning rate, and  $C$  represents the cost function. The choice of this function is related to the type of learning (such as supervised learning, unsupervised learning, and reinforcement learning) and the activation function.

DNN is currently the basis of many artificial intelligence applications. Due to the breakthrough applications of DNN in speech recognition and image recognition, the number of applications using DNN has exploded. These DNNs are deployed in various applications ranging from self-driving cars, cancer detection to complex games. In these many fields, DNN can surpass human accuracy. The outstanding performance of DNN comes from its ability to use statistical learning methods to extract high-level features from raw sensory data, and to obtain an effective representation of the input space from a large amount of data. This is different from previous methods that used manual feature extraction or expert design rules.

However, the price of DNN to obtain superior accuracy is high computational complexity cost. Although general-purpose computing engines (especially GPUs) have become the mainstay of DNN processing, it is becoming increasingly popular to provide specific acceleration methods for DNN computing.

According to different applications, the shape and size of deep neural networks are also different. Popular shapes and sizes are evolving rapidly to improve model accuracy and efficiency. The input of all deep neural networks is a set of information values that characterize the network to be analyzed and processed. These values can be the pixels of a picture, or the sample amplitude of a piece of audio, or a digital representation of a system or game state.

There are two main forms of input processing networks: feedforward and loop. In the feedforward network, all calculations are a series of operations based on the output of the previous layer. The final set of operations is the output of the network. In this type of deep neural network, the network has no memory, and the output is always independent of the input order of the previous network. In contrast, recurrent networks (LSTM is a popular variant) have internal memory, allowing long-term dependencies to affect the output. In these networks, the state values of some intermediate operations are stored in the network and are also used as inputs for other operations related to processing the latter input.

## 2.2 Transformer-based Language Model

Transformer is a sequence-to-sequence NLP model (Vaswani et al., 2017) which uses an at-

tention mechanism to draw global dependencies between input and output.

It takes a sequence of word embeddings from one vocabulary set as input and generates the probability of tokens in the other vocabulary set. The model mainly consists of encoding and decoding components. The encoding component is a stack of encoders that are all identical in structure but their weights are trained independently.

Likewise, the decoding component is also a stack of decoders. Note, the number of encoders and decoders can be adjusted to arrive at different transformer models. For instance, BERT only contains encoders (Lan et al., 2019). BERT<sub>BASE</sub> has 12 layers in the encoder stack while BERT<sub>LARGE</sub> has 24 layers in the encoder stack. OpenAI GPT-3 (Brown et al., 2020) only has 12 layers of decoders.

Before the word embedding is processed by the encoder, we add the position information of the tokens to the sequence so that the model is aware of the position of the token in each sequence. We use sine and cosine functions to encode this position information (Vaswani et al., 2017):

$$\mathbf{PE}_{(pos,2i)} = \sin(pos/10000^{2i/d_{model}}), \quad (2)$$

$$\mathbf{PE}_{(pos,2i+1)} = \cos(pos/10000^{2i/d_{model}}), \quad (3)$$

where  $pos$  is the position of the token in the sequence;  $d_{model}$  is the dimension of each word's embedding, and  $i$  is the  $i$ -th dimension in the word embedding vector. These positional values are added to the embedding of each token.

Self-attention is the key for an encoder. During linear transformation, the input  $\mathbf{X}$  is multiplied by three weight matrices, i.e., query ( $\mathbf{W}_Q$ ), key ( $\mathbf{W}_K$ ), and value ( $\mathbf{W}_V$ ), to arrive at  $\mathbf{Q}$ ,  $\mathbf{K}$ , and  $\mathbf{V}$ , respectively. That is  $\mathbf{Q} = \mathbf{X} \cdot \mathbf{W}_Q^T$ ,  $\mathbf{K} = \mathbf{X} \cdot \mathbf{W}_K^T$  and  $\mathbf{V} = \mathbf{X} \cdot \mathbf{W}_V^T$ . Afterwards, we follow Equation 4 to perform attention computation:

$$\mathbf{Z} = \text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q} \cdot \mathbf{K}^T}{\sqrt{d_k}}\right) \cdot \mathbf{V}. \quad (4)$$

Note, a mask is applied to exclude certain word-to-word interactions before softmax. One popular masking mechanism is to set the lower triangle part of the masking as 0 and the upper triangle part as negative infinity. When applied to  $\mathbf{Q} \cdot \mathbf{K}^T$ , this mask actually prevents the position information of later words from affecting the earlier ones. Multi-head attention extends Equation 4. The attention model possesses multiple sets of  $\mathbf{W}_Q$ ,  $\mathbf{W}_K$ ,  $\mathbf{W}_V$ , which

allows the model to jointly attend to information from different representation subspaces at different positions and each set of weights is a head. Since there are multiple  $\mathbf{Z}$ 's of Equation 4, multi-head attention combines them by multiplying  $\mathbf{Z}$  with a weight matrix  $\mathbf{W}_O$ .  $\sqrt{d_k}$  is the dot product's scaling factor, where  $d_k = \frac{d_{model}}{H}$ , and  $H$  is the number of heads.

After the self-attention module, there is a multi-layer perceptron (MLP) module consisting of two linear transformation layers with an activation layer between them. A layer normalization is applied after self-attention and MLP respectively, whose input will be added by the input of the module.

### 2.3 Network Pruning

As described in the background section, network pruning is an intuitive yet very effective strategy to reduce network size, thereby reduces the computation and time cost. Since there are countless researchers taking great effort on this, many pruning methods have been provided so far. For example, structure pruning, channel pruning, irregular pruning, and so on. Each algorithm has shown great performance. However, pruning usually causes a trade-off with an accuracy drop. Thus, our work becomes dealing with such trade-offs, more intuitively, pruning the network as much as possible and maintain the accuracy at the same time.

In our project, we have been working on a pruning method called pattern pruning. Pattern pruning prunes the whole row or column of the weight matrices. Therefore, such a pruning method benefits hardware even more. Nevertheless, pattern pruning also faces an accuracy drop as mentioned above. We also implement the ADMM algorithm to overcome the shortcoming.

ADMM is short for the alternating direction method of multipliers. The method is a regularized optimization during the DNN training, which can exactly enforce certain patterns in DNN weights (Yuan et al., 2021).

To be more explicit, we first generate masks with 0 and 1 in certain patterns. Then we apply masks to the weight matrix by multiplying them elementwise. Then the masks are also applied on the gradient matrix in the same way in each training step, thus the pruned weights do not update. The process is called hard-prune. One thing that's worth noting is that the algorithm also allows each weight matrix to select the best pattern from a pattern set

according to its L2 norm.

## 3 Experiment

In this section, we make a description of our experiment process on both computer vision model pruning and Transformer-based language model pruning.

### 3.1 Computer Vision Model Pruning

Reading and writing weights in memory is an energy-consuming and time-consuming process. Approximately, reading  $1.2MB$  data requires  $1second$  and  $1.47mj$  energy. Writing  $1.3MB$  data requires  $1second$  and  $1.51mj$  energy. It is more frequent while operating in energy-aware inference models since the models are needed to be switched based on the energy level of the sensing devices. Therefore, technically, if the models of different sparsities have shared weights, the erasing and rewriting of the shared weights can be safely skipped, thus, the switching process could be more energy efficient.

Normally, model compressing process contains a retrain step to guarantee the mask is applied on the weight matrix. However, the retraining process naturally changes the weights remained while applying masks, thus shared weights by implementing conventional weight pruning strategy is not possible. Alternatively, we simplified the process through removing the retraining step. By sacrificing an acceptable accuracy drop, the shared weights among models of different sparsities become possible.

In our shared-weights training (for demonstration, we use three models), the weights of the pruned layer are:

- *HighSparsity* :  $W_h = W_{Initial} \times M_h$
- *MediumSparsity* :  $W_m = W_{Initial} \times M_{hm} + W_h$
- *LowSparsity* :  $W_l = W_{Initial} \times M_{hm} + W_m$

In the equations shown above,  $W$  denotes the weight matrix,  $M$  represents the mask. The *footnote* stands for the percentage of sparsity.  $h$ ,  $m$ ,  $l$  and  $hm$  stand for *high*, *medium*, *low* and *slightly higher than medium* sparsity, respectively. It will be illustrated further in the Fig. 1, in which the shared-weights training workflow is illustrated.




Figure 1: Shared-Weights Training Workflow

The white cells in the figure represent 0s. The grey cells represent 1s. Orange, green and blue cells are remained weights after pruning. The cells illustrated in the same color stand for shared weights.

The algorithm is simple yet effective, however, additional indices are required to store the positions of pruned weights. To compensate the memory spent on storing indices, pruning ratios have to be larger than a certain threshold. For the implementation of the LeNet5 model on CPU based devices, the sparsity needs to be no less than 4% to compensate the storage of indices, and the threshold is 20% when implementing on LEA based devices.

### 3.2 Transformer-based Language Model Pruning

In this work, various weight pruning methods, such as irregular pruning, column pruning, row pruning and whole block pruning will be explored on  $BERT_{BASE}$  to demonstrate the effectiveness of weight pruning in shrinking model size, accelerating model inference time while retaining model performance.

We mainly focus on the backbone pruning of  $BERT_{BASE}$ , which means we prune the 6 weight matrices (i.e., query, key, value, attention.output.dense, intermediate.dense, output.dense) on 12 encoders. We follow the identical layer rule for pruning ratio, which implies we apply the same pruning ratio for each layer on each encoder. For each pruning method, we first fine-tune the pre-trained model on corresponding tasks. Then, we prune the model with the pruning method. Finally, we retrain the model until it's converged.

For irregular pruning, we eliminate the least absolute value of weights (Han et al., 2015b) after finetuning. For column pruning, we calculated the l2 norm of each column of a specific weight and fill the least l2 norm value corresponding column values with zeros. Row pruning and whole block

pruning follow the same way with different pattern.

#### 3.2.1 Dataset

We conduct experiments on GLUE benchmark (Wang et al., 2018), a comprehensive collection of natural language understanding tasks covering three NLP categories, i.e., inference tasks (MNLI (Williams et al., 2018), Quora Question Pairs (QQP) (Zhang et al., 2018b), QNLI (Wang et al., 2018) (a set of over 100,000+ question-answer pairs from SQuAD (Rajpurkar et al., 2016)), single-sentence (SST-2 (Socher et al., 2013)), paraphrase similarity matching (STS-B (Cer et al., 2017), Microsoft Research Paraphrase Corpus (MRPC (Dolan and Brockett, 2005))).

#### 3.2.2 Baseline Models

The baseline model is our own fine-tuned unpruned  $BERT_{BASE}$  (Devlin et al., 2019). We report our results (from the official bert-base-uncased) as Full  $BERT_{BASE}$ . We use Huggingface Transformer toolkit (Wolf et al., 2019) to conduct our experiment. There are 12 layers ( $L=12$ ; hidden size  $H=768$ ; self-attention heads  $A=12$ ), with 110 million parameters.

#### 3.2.3 Evaluation Metrics

We apply the same metrics of the tasks as the GLUE paper (Wang et al., 2018), i.e., accuracy scores are reported for RTE and QNLI; F1 scores are reported for MRPC; Spearman correlations are reported for STS-B.

#### 3.2.4 Platforms

We use Python 3.6.10 with PyTorch 1.4.0 and CUDA 11.1 on Quadro RTX6000 GPU and Intel(R) Xeon(R) Gold 6244 @ 3.60GHz CPU.

#### 3.2.5 Experimental Results

- CV pruning: Table 2 shows the accuracy of different sparsities after shared-weights training. We can tell from the table that shared-weights training can help retain weights among different sparsities to have them shared with a comparable accuracy. The patterns applied on the weight matrix are selected manually according to the rule proposed in the PCONV paper (Ma and et.al., 2020) and the highest accuracy is given. As for *LeNet5*, the high sparsity is 60.24% and the corresponding accuracy is 98.94%. The medium Sparsity is 20.093% and its accuracy is 98.96%. The low sparsity is 8.47% and its accuracy is



Models	MNLI	QQP	QNLI	SST-2	STS-B	MRPC	RTE
<b>BERT<sub>BASE</sub> (Devlin et al., 2018)</b>	84.6	91.2	90.5	93.5	85.8	88.9	66.4
<b>BERT<sub>BASE</sub> (ours)</b>	83.9	91.4	91.1	92.7	85.8	89.8	66.4
<b>BERT<sub>BASE</sub> irregular prune (ours)</b>	81.1	86.3	88.9	89.3	83.1	89.3	66.1
<b>BERT<sub>BASE</sub> row prune (ours)</b>	78.1	86.8	88.1	85.5	78.7	77.7	53.2
<b>BERT<sub>BASE</sub> column prune (ours)</b>	78.5	87	88.3	85.9	78.7	77.8	53.8
<b>BERT<sub>BASE</sub> whole block prune (ours)</b>	78.8	87.3	86.2	87.3	85.3	79.2	53.1
<b>Compression rate</b>	1.428×	1.428×	1.428×	1.428×	1.428×	1.428×	1.428×

Table 1: Comparison of test accuracy using BERT<sub>BASE</sub> among the nine GLUE benchmark tasks.

Dataset/Task		MNIST (Image Classification)			HAR (Human Activity Recognition)			OKG (Speech Recognition)			ImageNet (Image Classification)		
Models		LeNet-5 (LeCun et al., 1998)			HAR-Net(Anguita and et al., 2013)			OKG-Net(Warden, 2018)			SqNxt-23 (Gholami and et. al., 2018)		
		M1	M2	M3	M1	M2	M3	M1	M2	M3	M1	M2	M3
SW-Train	Sparsity	60.24%	40.50%	20.70%	24.13%	5.31%	2.69%	83.76%	56.32%	30.78%	55.39%	54.30%	54.07%
	Accuracy	98.94%	98.96%	99.02%	90.19%	89.78%	87.21%	82.01%	83.43%	83.5%	73.80%	75.02%	77.36%
	Latency (s)	1.1	1.3	1.5	0.66	0.83	0.85	0.82	1.3	1.62	3.4	3.6	3.7

Table 2: Results of Shared Weight Training

99.02%. Part of the remaining weights are shared, as illustrated in Figure 1. Similarly in the *SqNxt* – 23 model, the accuracy slightly increases as the sparsity reduces. Apart from image classification tasks, our shared-weights training shows satisfying accuracy on recognition tasks as well. For example, as for the *HAR* – *Net* model, the high, medium and low sparsities are 24.13%, 5.31% and 2.69%, respectively and their corresponding accuracy are 90.19%, 89.78% and 87.21%. Comparably, the accuracy and sparsity of the OKG-Net show the same trend as *HAR* – *Net*. Thus, with sacrificing little accuracy drop, the shared-weights training algorithm could keep weights unchanged, and further benefit the inference process on energy harvesting devices.

## 4 Conclusion

In order to tackle the problem of accommodating DNN models on hardware, pruning is a key point. This can usually causes a trade-off with an accuracy drop. In this project, we focus on dealing with such trade-offs, more intuitively, pruning the network and maintain the accuracy at the same time. The experiments are based on the computer vision model pruning and Transformer-based language model pruning. As for the results, we can make a conclusion that pruning strategy has benefits on multiple-tasks in both hardware and software level, by accelerating inference process and saving computation costs.

## References

- Davide Anguita and et.al. 2013. A public domain dataset for human activity recognition using smart-phones. In *Esann*, volume 3, page 3.
- Sajid Anwar, Kyuyeon Hwang, and Wonyong Sung. 2017. Structured pruning of deep convolutional neural networks. *ACM Journal on Emerging Technologies in Computing Systems (JETC)*, 13(3):1–18.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, Chris Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. Language models are few-shot learners. In *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901, Virtual. Curran Associates, Inc.
- Daniel Cer, Mona Diab, Eneko Agirre, Iñigo Lopez-Gazpio, and Lucia Specia. 2017. SemEval-2017 task 1: Semantic textual similarity multilingual and crosslingual focused evaluation. In *Proceedings of the 11th International Workshop on Semantic Evaluation (SemEval-2017)*, pages 1–14, Vancouver, Canada. The Association for Computer Linguistics.
- Xue Chen, Lanyong Zhang, Tong Liu, and MM Kamruzzaman. 2019. Research on deep learning in the field of mechanical equipment fault diagnosis image quality. *Journal of Visual Communication and Image Representation*, 62:402–409.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2018. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*.

430	Jacob Devlin et al. 2019. Bert: Pre-training of deep	484
431	bidirectional transformers for language understand-	485
432	ing. In <i>NAACL-HLT (1)</i> .	486
433	Bill Dolan and Chris Brockett. 2005. Automati-	487
434	cally constructing a corpus of sentential paraphrases.	488
435	In <i>Third International Workshop on Paraphrasing</i> .	489
436	Asia Federation of Natural Language Processing.	
437	Colin Garvey. 2018. A framework for evaluating bar-	490
438	riers to the democratization of artificial intelligence.	491
439	In <i>Thirty-Second AAAI Conference on Artificial In-</i>	492
440	<i>telligence</i> .	493
441	Amir Gholami and et. al. 2018. Squeezenext:	494
442	Hardware-aware neural network design. In <i>CVPR</i> ,	495
443	pages 1638–1647.	496
444	Song Han, Huizi Mao, and William J Dally. 2015a.	497
445	Deep compression: Compressing deep neural net-	
446	works with pruning, trained quantization and huff-	498
447	man coding. <i>arXiv preprint arXiv:1510.00149</i> .	499
448	Song Han et al. 2015b. Learning both weights and con-	500
449	nections for efficient neural network. In <i>Advances in</i>	501
450	<i>neural information processing systems</i> , pages 1135–	
451	1143.	502
452	Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. 2015.	503
453	Distilling the knowledge in a neural network. <i>arXiv</i>	504
454	<i>preprint arXiv:1503.02531</i> .	
455	Zhenzhong Lan, Mingda Chen, Sebastian Goodman,	505
456	Kevin Gimpel, Piyush Sharma, and Radu Soricut.	506
457	2019. <i>Albert: A lite bert for self-supervised learn-</i>	507
458	<i>ing of language representations</i> .	
459	Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick	508
460	Haffner. 1998. Gradient-based learning applied to	509
461	document recognition. <i>Proceedings of the IEEE</i> ,	510
462	86(11):2278–2324.	511
463	Christos Louizos, Max Welling, and Diederik P	512
464	Kingma. 2017. Learning sparse neural net-	513
465	works through $l_0$ regularization. <i>arXiv preprint</i>	514
466	<i>arXiv:1712.01312</i> .	515
467	Xiaolong Ma and et.al. 2020. Pconv: The missing	516
468	but desirable sparsity in dnn weight pruning for real-	
469	time execution on mobile devices. <i>ArXiv</i> .	517
470	Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and	518
471	Percy Liang. 2016. Squad: 100,000+ questions for	519
472	machine comprehension of text. In <i>Proceedings of</i>	
473	<i>the 2016 Conference on Empirical Methods in Natu-</i>	520
474	<i>ral Language Processing</i> , pages 2383–2392, Austin,	521
475	Texas. Association for Computational Linguistics.	522
476	Richard Socher, Alex Perelygin, Jean Wu, Jason	523
477	Chuang, Christopher D. Manning, Andrew Ng, and	524
478	Christopher Potts. 2013. Recursive deep models	525
479	for semantic compositionality over a sentiment tree-	526
480	bank. In <i>Proceedings of the Conference on Empiri-</i>	527
481	<i>cal Methods in Natural Language Processing</i> , pages	528
482	1631–1642, Seattle, Washington, USA. Association	529
483	for Computational Linguistics.	530
	Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob	531
	Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz	532
	Kaiser, and Illia Polosukhin. 2017. Attention is all	533
	you need. In <i>Advances in neural information pro-</i>	534
	<i>cessing systems</i> , pages 5998–6008, Red Hook, NY,	535
	USA. Curran Associates Inc.	536
	Alex Wang, Amanpreet Singh, Julian Michael, Felix	537
	Hill, Omer Levy, and Samuel Bowman. 2018. Glue:	538
	A multi-task benchmark and analysis platform for	539
	natural language understanding. In <i>Proceedings</i>	540
	<i>of the 2018 EMNLP Workshop BlackboxNLP: An-</i>	
	<i>alyzing and Interpreting Neural Networks for NLP</i> ,	
	pages 353–355, Brussels, Belgium. Association for	
	Computational Linguistics.	
	Yijue Wang, Chenghong Wang, Zigeng Wang,	
	Shanglin Zhou, Hang Liu, Jinbo Bi, Caiwen Ding,	
	and Sanguthevar Rajasekaran. 2021. <i>Against mem-</i>	
	<i>bership inference attack: Pruning is all you need</i> .	
	Ziheng Wang, Jeremy Wohlwend, and Tao Lei. 2019.	
	Structured pruning of large language models. <i>arXiv</i>	
	<i>preprint arXiv:1910.04732</i> .	
	Pete Warden. 2018. Speech commands: A dataset	
	for limited-vocabulary speech recognition. <i>ArXiv</i> ,	
	abs/1804.03209.	
	Adina Williams, Nikita Nangia, and Samuel Bowman.	
	2018. A broad-coverage challenge corpus for sen-	
	tence understanding through inference. In <i>Proceed-</i>	
	<i>ings of the 2018 Conference of the North American</i>	
	<i>Chapter of the Association for Computational Lin-</i>	
	<i>guistics: Human Language Technologies, Volume</i>	
	<i>1 (Long Papers)</i> , pages 1112–1122, New Orleans,	
	Louisiana. Association for Computational Linguis-	
	tics.	
	Thomas Wolf et al. 2019. Huggingface’s transformers:	
	State-of-the-art natural language processing. <i>ArXiv</i> ,	
	abs/1910.03771.	
	Xiyu Yu, Tongliang Liu, Xinchao Wang, and Dacheng	
	Tao. 2017. On compressing deep models by low	
	rank and sparse decomposition. In <i>Proceedings of</i>	
	<i>the IEEE Conference on Computer Vision and Pat-</i>	
	<i>tern Recognition</i> , pages 7370–7379.	
	Geng Yuan, Payman Behnam, Zhengang Li, Ali	
	Shafiee, Sheng Lin, Xiaolong Ma, Hang Liu, Xue-	
	hai Qian, Mahdi Nazm Bojnordi, Yanzhi Wang, and	
	Caiwen Ding. 2021. <i>Forms: Fine-grained polarized</i>	
	<i>reram-based in-situ computation for mixed-signal</i>	
	<i>dnn accelerator</i> .	
	Tianyun Zhang, Shaokai Ye, Kaiqi Zhang, Jian Tang,	
	Wujie Wen, Makan Fardad, and Yanzhi Wang.	
	2018a. A systematic dnn weight pruning framework	
	using alternating direction method of multipliers. In	
	<i>Proceedings of the European Conference on Com-</i>	
	<i>puter Vision (ECCV)</i> , pages 184–199.	
	Wei Zhang, Lu Hou, Yichun Yin, Lifeng Shang, Xiao	
	Chen, Xin Jiang, and Qun Liu. 2020. Ternarybert:	
	Distillation-aware ultra-low bit bert. <i>arXiv preprint</i>	
	<i>arXiv:2009.12812</i> .	

541 Xiaodong Zhang, Xu Sun, and Houfeng Wang. 2018b.  
542 Duplicate question identification by integrating  
543 framenet with neural networks. In *Proceedings of*  
544 *the Conference on Artificial Intelligence*, volume 32,  
545 New Orleans, Louisiana, USA. AAAI Press.